

CSE475 (Fall 2025) Mini-Project Brief: Self-Supervised Learning on Provided Datasets

Department of Computer Science and Engineering
East West University, Dhaka, Bangladesh
Primary execution environment: Kaggle (required)

Abstract

This document specifies the eight-week mini-project for CSE475 (Fall 2025), focused on applying three self-supervised learning (SSL) methods to a provided dataset (or dataset pair). It defines task objectives, deliverables, evaluation metrics, checkpointing/resumability requirements, marking scheme, and reporting standards. All work must be reproducible on Kaggle.

1. Goal and Setting

Over eight weeks, student teams will execute a complete study applying three SSL methods to a provided dataset (or dataset pair). The work must be fully reproducible on Kaggle; Colab or local experimentation is optional for comparison only. Submit Kaggle notebooks, a paper-style report, slides, and saved model artifacts.

2. Task 1 — EDA and Related Work (4.0 marks)

2.1. Image-focused EDA (with concise figures/tables)

Report: (i) RGB and HSV histograms; per-image and per-class mean/standard deviation; brightness/contrast spread; saturation clipping; (ii) resolution and aspect-ratio distributions; resizing/padding strategy and class-wise effects; (iii) sharpness/noise using Laplacian variance and a light noise proxy; (iv) white-balance sanity via a gray-world check; (v) duplicate detection using perceptual hashing; (vi) class balance; (vii) leakage safeguards in splitting (e.g., grouping by source/ID). Run a small augmentation probe (crop/flip/color jitter/blur) to identify safe vs. harmful augmentations.

2.2. Related Work (dataset/domain)

Create a *Related Work Summary* focused on the same dataset or a closely related domain. This table *constrains Task 2*: supervised backbones must be drawn from credible related work on this dataset/domain.

Required Table Columns (exact). *Title — Dataset name and URL — Dataset description (samples, classes, images per class or per split) — Methods name — Accuracy of the model — Pros — Cons — Citation.*

A LaTeX template is provided in Section 10.

Deliverables: one Kaggle EDA notebook; 1–2 page related-work summary with the required table and citations.

3. Task 2 — Supervised Baselines from Related Work (5.0 marks)

Implement at least five pretrained supervised backbones *explicitly* drawn from models reported in the Task-1 Related Work table (e.g., GoogLeNet, MobileNet, ResNet-50/101, EfficientNet, ViT—only if the related work in this domain used them).

Splits: evaluate across {90:10, 80:20, 70:30, 60:40, 50:50, 40:60, 30:70, 20:80, 10:90} (test set fixed per ratio). Within each training portion, allocate 10% for validation. Training: train until a clear learning-curve plateau or for at least 50 epochs; show training/validation curves.

Report (per model and split): Accuracy, per-class Accuracy, Precision, Recall, F1-score, ROC–AUC with ROC curves, Confusion Matrix, training/testing wall-clock time, and GFLOPs per inference at a fixed input size (state the size). Summarize failure modes (confused class pairs). Identify the best supervised baseline as the reference for SSL.

Deliverables: Kaggle notebook(s) with consistent logs; a brief table highlighting the top-1 supervised result.

4. Task 3 — Select and Explain Three SSL Methods (2.5 marks)

Choose three SSL methods suitable for your data (e.g., SimCLR, BYOL, Barlow Twins, DINO/DINOv2, iBOT, MAE, VICReg, SimSiam, SwaV). Justify each selection based on dataset characteristics. Prepare a teaching-style explanation (as if to a class-10 student): intuition, objective (with a friendly diagram), how positives/negatives or mask ing/tokenization work, and why the method fits the dataset.

Deliverables: 6–8 explanatory slides (or a narrative section) with citations.

5. Tasks 4–6 — Implement SSL-1, SSL-2, SSL-3 (3.0 marks each = 9.0)

5.1. Label-free Pretraining

Run the SSL pretext on the training split with a documented augmentation recipe. Save the frozen encoder.

5.2. Downstream Evaluation

- Linear probe on frozen features (primary SSL yardstick).
- Shallow heads: MLP, SVM, Decision Tree, Random Forest on frozen features.

Full fine-tuning of the encoder; compare to linear probe to quantify gains/costs.

5.3. Embedding Analysis

Provide t-SNE, UMAP, and PCA plots colored by true labels; report the Silhouette score to quantify separation.

5.4. Metrics & Logs

Accuracy, per-class Accuracy, Precision, Recall, F1, ROC–AUC with curves, Confusion Matrix, learning curves (pretraining proxy + downstream), train/test time, and GFLOPs (same input size across methods). Add k-NN accuracy ($k \in \{1, 5, 20\}$) in embedding space and label-efficiency curves (linear-probe accuracy at $\{1\%, 5\%, 10\%, 25\%, 50\%\}$ labeled data).

Deliverables: one Kaggle notebook per SSL method; a brief comparison table.

6. Task 7 — Ratios, Ablations, and Statistics (3.5 marks)

Re-run downstream evaluations for all SSLs across the same train:test grid used in Task 2. Perform at least one solid ablation per SSL (e.g., remove a key augmentation; change projection-head depth; vary batch size or pretraining epochs; alter optimizer/LR). Summarize effects on linear-probe and label-efficiency outcomes.

Statistics: use paired t-tests for two-model comparisons on identical folds; for three or more models (or multiple datasets), use Wilcoxon/Friedman with Nemenyi post-hoc when appropriate. Report p-values and discuss practical vs. statistical significance. Deliverables: a consolidated results notebook with ablation tables, statistical-test outputs, and 1–2 figures showing effect sizes.

7. Task 8 — Paper-Style Report and Final Presentation (4.0 marks)

Prepare an 8–12 page report: abstract, introduction, related work, methods (supervised baselines + three SSL methods), experimental setup, results (including embeddings and label-efficiency), ablations with statistics, discussion, limitations/ethics, and conclusion. Create 10–12 slides for a concise, defensible

presentation.

Deliverables: PDF report, slides (PDF/PPTX), public Kaggle project with “Run All” notebooks, and a mirroring GitHub repo (requirements, fixed seeds, README).

3

8. Model Saving, Checkpointing, and Resumability (Mandatory)

8.1. Where to Save

Save checkpoints to `/kaggle/working/` during training. At notebook end, persist artifacts by attaching them as notebook Outputs (so they remain available on rerun and can be converted into a Kaggle Dataset).

8.2. What to Save (Supervised and SSL)

Save both the latest and the best checkpoint (based on your validation metric). Each checkpoint must include:

- model state dict (include encoder and any projection/head modules),
- optimizer state dict, scheduler state dict (if used),
- epoch, global step, best metric (name and value),
- random states (Python, NumPy, PyTorch CPU/GPU),
- data split manifest (file list or seed + stratification info),
- method-specific buffers (e.g., EMA/momentum teacher for BYOL/DINO; queues/memory banks for MoCo/SwaV; tokenizers/positional embeddings for ViTs).

8.3. Best-Model Criterion

Supervised: best validation ROC–AUC (or top-1 accuracy when class imbalance is negligible—state and justify).

SSL: best *linear-probe* validation accuracy (use a consistent labeling budget across methods).

8.4. Resumable Training Loop (SSL required; supervised recommended)

Training must support *resume-from-checkpoint* restoring model, optimizer, scheduler, EMA/momentum teacher (if any), queues/buffers, epoch/step, RNG states, and logging offsets. Implement periodic and on-improvement saves; use early stopping with patience and save the best checkpoint. On resume, continue seamlessly and keep updating “latest” and “best”.

8.5. Determinism and Logging

Fix seeds; enable deterministic flags when compatible. Record library versions and CUDA/cuDNN info. Log metrics and hyperparameters to CSV/JSONL and save along side checkpoints.

4

9. Total Marks (25)

- Task 1: 4.0
- Task 2: 5.0
- Task 3: 2.5
- Task 4 (SSL-1): 3.0
- Task 5 (SSL-2): 3.0
- Task 6 (SSL-3): 3.0
- Task 7: 3.5
- Task 8: 4.0

10. Related Work Summary Table Template *Use this table in*

Task 1. Add/expand rows as needed. Keep entries concise.

Title and URL	Dataset name and description (same as in paper)	Dataset description (same as in paper)	Classes, images, methods	Accuracy of the model	Pros	Cons	Citation
	Study	$N=25,000$		$C=10$ classes; ~50, ViT B/16		simple recipe imbalanced; small images	
	XYZ images/class or per split)	ResNet	Top-1: Strong [?, ?] on XYZ Images (https://example.com/images;	92.3 %; AUC: 0.964		generalization;	
Example		Data					
	com/ xyz)	; official 2,500/class/train/val/test					

References

1. T. Fawcett, “An introduction to ROC analysis,” *Pattern Recognition Letters*, vol. 27, no. 8, pp. 861–874, 2006. doi: 10.1016/j.patrec.2005.10.010.
2. J. Demšar, “Statistical comparisons of classifiers over multiple data sets,” *Journal of Machine Learning Research*, vol. 7, pp. 1–30, 2006.
3. L. van der Maaten and G. Hinton, “Visualizing data using t-SNE,” *JMLR*, vol. 9, pp. 2579–2605, 2008.
4. L. McInnes, J. Healy, and J. Melville, “UMAP: Uniform Manifold Approximation and Projection for Dimension Reduction,” arXiv:1802.03426, 2018.
5. P. J. Rousseeuw, “Silhouettes: A graphical aid to the interpretation and validation of cluster analysis,” *Journal of Computational and Applied Mathematics*, vol. 20, pp. 53–65, 1987.
6. T. Chen, S. Kornblith, M. Norouzi, and G. Hinton, “A simple framework for contrastive learning of visual representations,” in *Proc. ICML*, 2020. arXiv:2002.05709.
7. J.-B. Grill *et al.*, “Bootstrap your own latent: A new approach to self-supervised learning,” in *Proc. NeurIPS*, 2020. arXiv:2006.07733.
8. J. Zbontar *et al.*, “Barlow twins: Self-supervised learning via redundancy reduction,” in *Proc. ICML*, 2021. arXiv:2103.03230.
9. M. Caron *et al.*, “Emerging properties in self-supervised vision transformers,” in *Proc. ICCV*, 2021. arXiv:2104.14294.
10. M. Oquab *et al.*, “DINOv2: Learning robust visual features without supervision,” 2023. arXiv:2304.07193.
11. K. He, X. Chen, S. Xie, Y. Li, P. Dollár, and R. Girshick, “Masked autoencoders are scalable vision learners,” in *Proc. CVPR*, 2022. arXiv:2111.06377.
12. K. Zhou *et al.*, “iBOT: Image BERT pre-training with online tokenizer,” in *Proc. ICLR*, 2022. arXiv:2111.07832.

13. A. Bardes, J. Ponce, and Y. LeCun, “VICReg: Variance-Invariance-Covariance regularization for self-supervised learning,” in *Proc. ICLR*, 2022. arXiv:2105.04906.